

Assignment 3 : IBM-Project-10469-1664355350

Problem Statement: Abalone Age Prediction

Description :

Predicting the age of abalone from physical measurements. The age of abalone is determined by cutting the shell through the cone, staining it, and counting the number of rings through a microscope -- a boring and time-consuming task. Other measurements, which are easier to obtain, are used to predict age. Further information, such as weather patterns and location (hence food availability) may be required to solve the problem.

Importing Modules

In [1]:

```
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
import numpy as np
```

1. Dataset has been downloaded

In []:

```
#Name of the dataset: abalone.csv
```

2. Load the dataset into the tool

In [12]:

```
data=pd.read_csv("abalone.csv")
data.head()
```

Out[12]:

	Sex	Length	Diameter	Height	Whole weight	Shucked weight	Viscera weight	Shell weight	Rings
0	M	0.455	0.365	0.095	0.5140	0.2245	0.1010	0.150	15
1	M	0.350	0.265	0.090	0.2255	0.0995	0.0485	0.070	7
2	F	0.530	0.420	0.135	0.6770	0.2565	0.1415	0.210	9
3	M	0.440	0.365	0.125	0.5160	0.2155	0.1140	0.155	10
4	I	0.330	0.255	0.080	0.2050	0.0895	0.0395	0.055	7

Let's know the shape of the data

In [13]:

```
data.shape
```

Out[13]:

```
(4177, 9)
```

One additional task is that, we have to add the "Age" column using "Rings" data. We just have to add '1.5' to the ring data

In [14]:

```
Age=1.5+data.Rings
data["Age"]=Age
data=data.rename(columns = {'Whole_weight':'Whole_weight', 'Shucked_weight': 'Shucked_weight', 'Viscera_weight': 'Viscera_weight',
                             'Shell_weight': 'Shell_weight'})
data=data.drop(columns=["Rings"],axis=1)
data.head()
```

Out[14]:

	Sex	Length	Diameter	Height	Whole_weight	Shucked_weight	Viscera_weight	Shell_weight	Age
0	M	0.455	0.365	0.095	0.5140	0.2245	0.1010	0.150	16.5
1	M	0.350	0.265	0.090	0.2255	0.0995	0.0485	0.070	8.5
2	F	0.530	0.420	0.135	0.6770	0.2565	0.1415	0.210	10.5
3	M	0.440	0.365	0.125	0.5160	0.2155	0.1140	0.155	11.5
4	I	0.330	0.255	0.080	0.2050	0.0895	0.0395	0.055	8.5

3. Perform Below Visualizations.

(i) Univariate Analysis

#

The term univariate analysis refers to the analysis of one variable. You can remember this because the prefix “uni” means “one.” There are three common ways to perform univariate analysis on one variable: 1. Summary statistics – Measures the center and spread of values.

#


Histogram


In [16]:

```
sns.displot(data["Age"], color='darkorange')
```

Out[16]:

<seaborn.axisgrid.FacetGrid at 0x7fd3f837a430>






In [103]:

```
sns.histplot(y=data.Age,color='pink')
```

Out[103]:

<AxesSubplot:xlabel='Count', ylabel='Age'>




In [106]:

```
sns.histplot(x=data.Age,color='green')
```

Out[106]:

<AxesSubplot:xlabel='Age', ylabel='Count'>




Boxplot


In [52]:

```
sns.boxplot(x=data.Age,color='pink')
```

Out[52]:

<AxesSubplot:xlabel='Age'>






Countplot

In [51]:

```
sns.countplot(x=data.Age)
```

Out[51]:

<AxesSubplot:xlabel='Age', ylabel='count'>



(ii) Bi-Variate Analysis

#

Image result for bivariate analysis in python It is a methodical statistical technique applied to a pair of variables (features/ attributes) of data to determine the empirical relationship between them. In order words, it is meant to determine any concurrent relations (usually over and above a simple correlation analysis).

#


Barplot


In [50]:

```
sns.barplot(x=data.Height,y=data.Age)
```

Out[50]:

<AxesSubplot:xlabel='Height', ylabel='Age'>






Linearplot

In [49]:

```
sns.lineplot(x=data.Age,y=data.Height, color='purple')
```

Out[49]:

<AxesSubplot:xlabel='Age', ylabel='Height'>




Scatterplot

In [42]:

```
sns.scatterplot(x=data.Age,y=data.Height,color='green')
```

Out[42]:

<AxesSubplot:xlabel='Age', ylabel='Height'>




Pointplot

In [45]:

```
sns.pointplot(x=data.Age, y=data.Height, color="blue")
```

Out[45]:

<AxesSubplot:xlabel='Age', ylabel='Height'>




Regplot

In [48]:

```
sns.regplot(x=data.Age,y=data.Height,color='orange')
```

Out[48]:

```
<AxesSubplot:xlabel='Age', ylabel='Height'>
```



(iii) Multi-Variate Analysis

#

Multivariate analysis is based in observation and analysis of more than one statistical outcome variable at a time. In design and analysis, the technique is used to perform trade studies across multiple dimensions while taking into account the effects of all variables on the responses of interest.

#

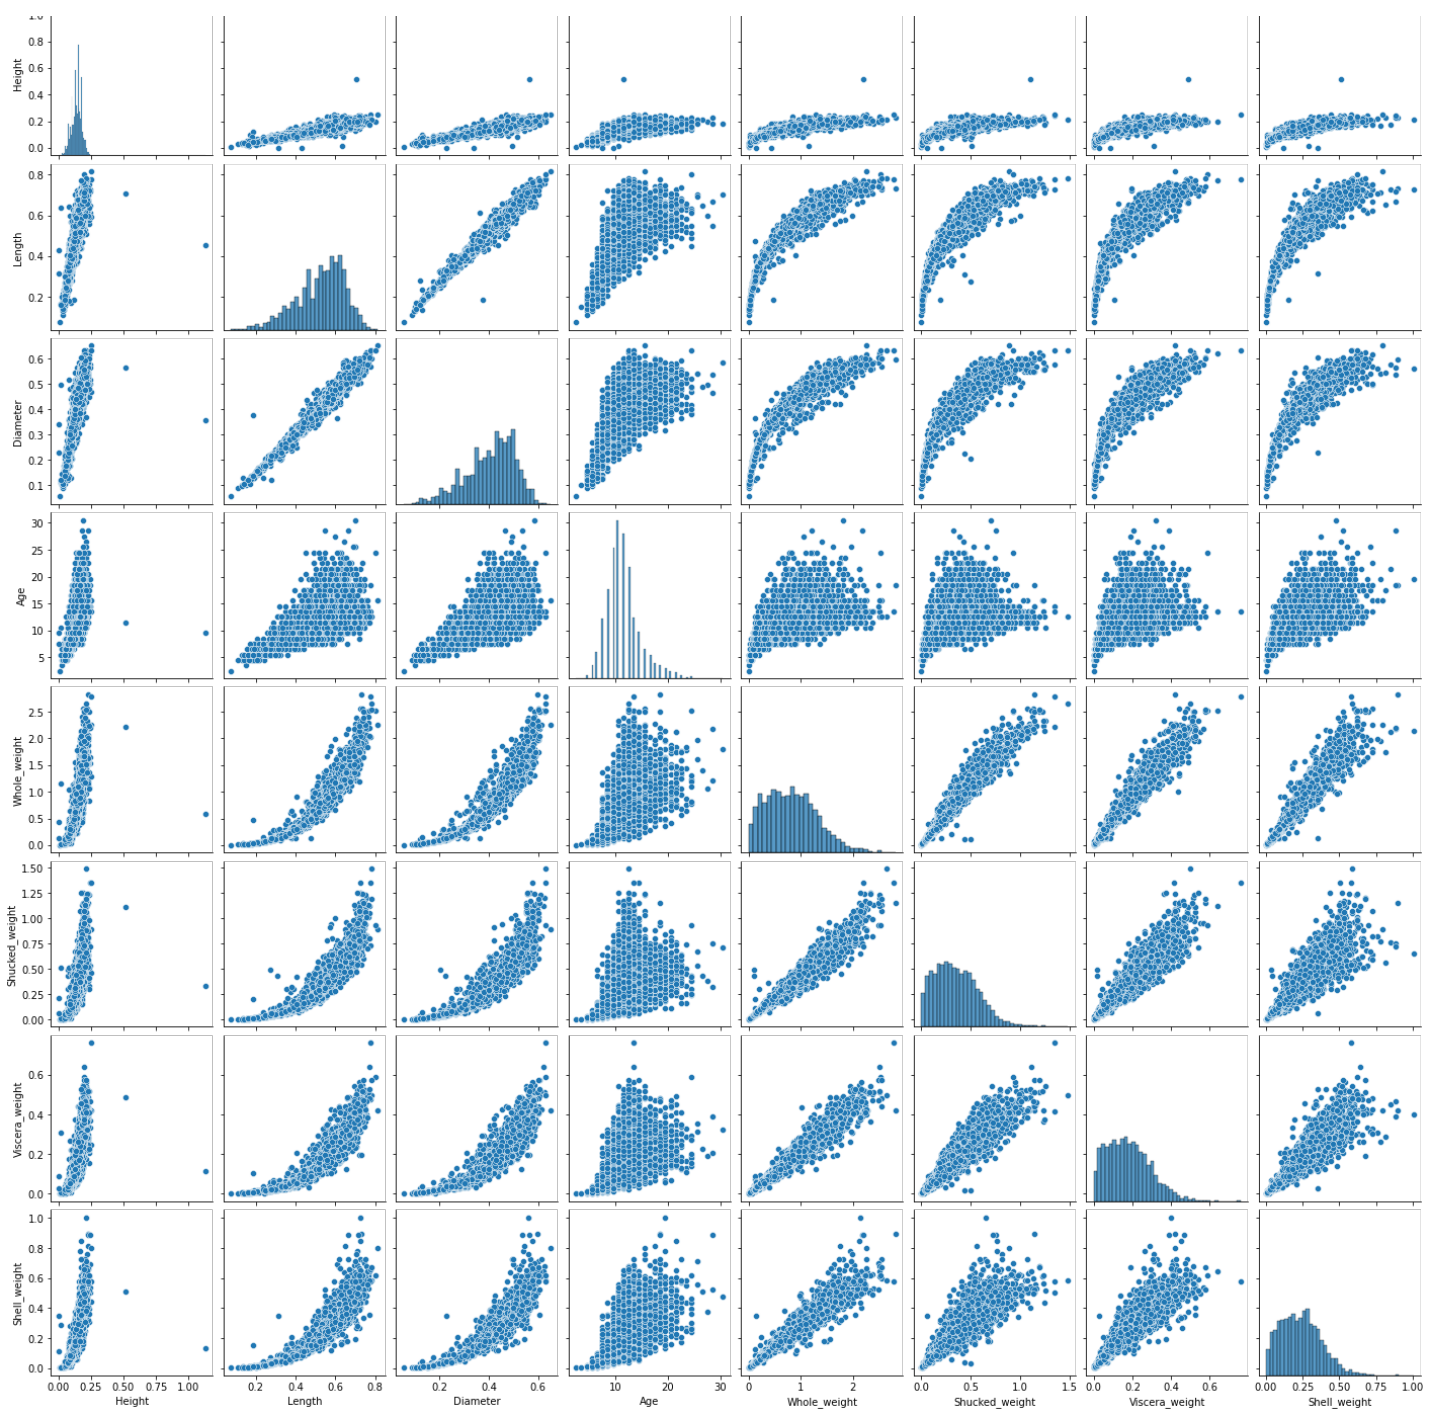
Pairplot

In [57]:

```
sns.pairplot(data=data[["Height", "Length", "Diameter", "Age", "Whole_weight", "Shucked_weight", "Viscera_weight", "Shell_weight"]])
```

Out[57]:

```
<seaborn.axisgrid.PairGrid at 0x7fd3d93e1040>
```

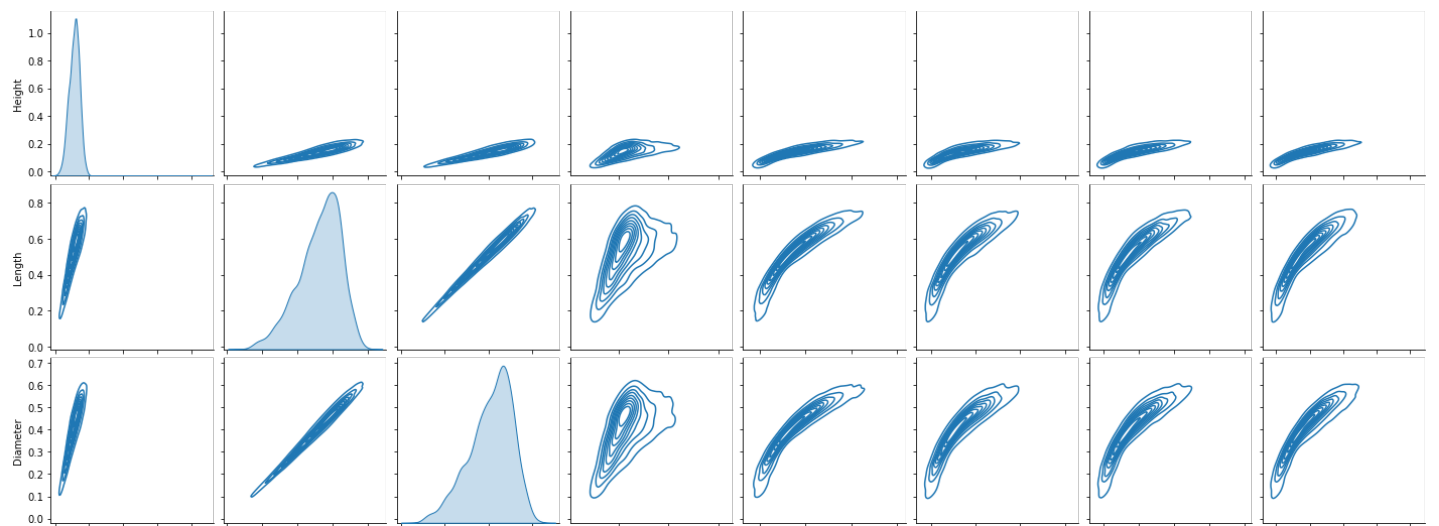


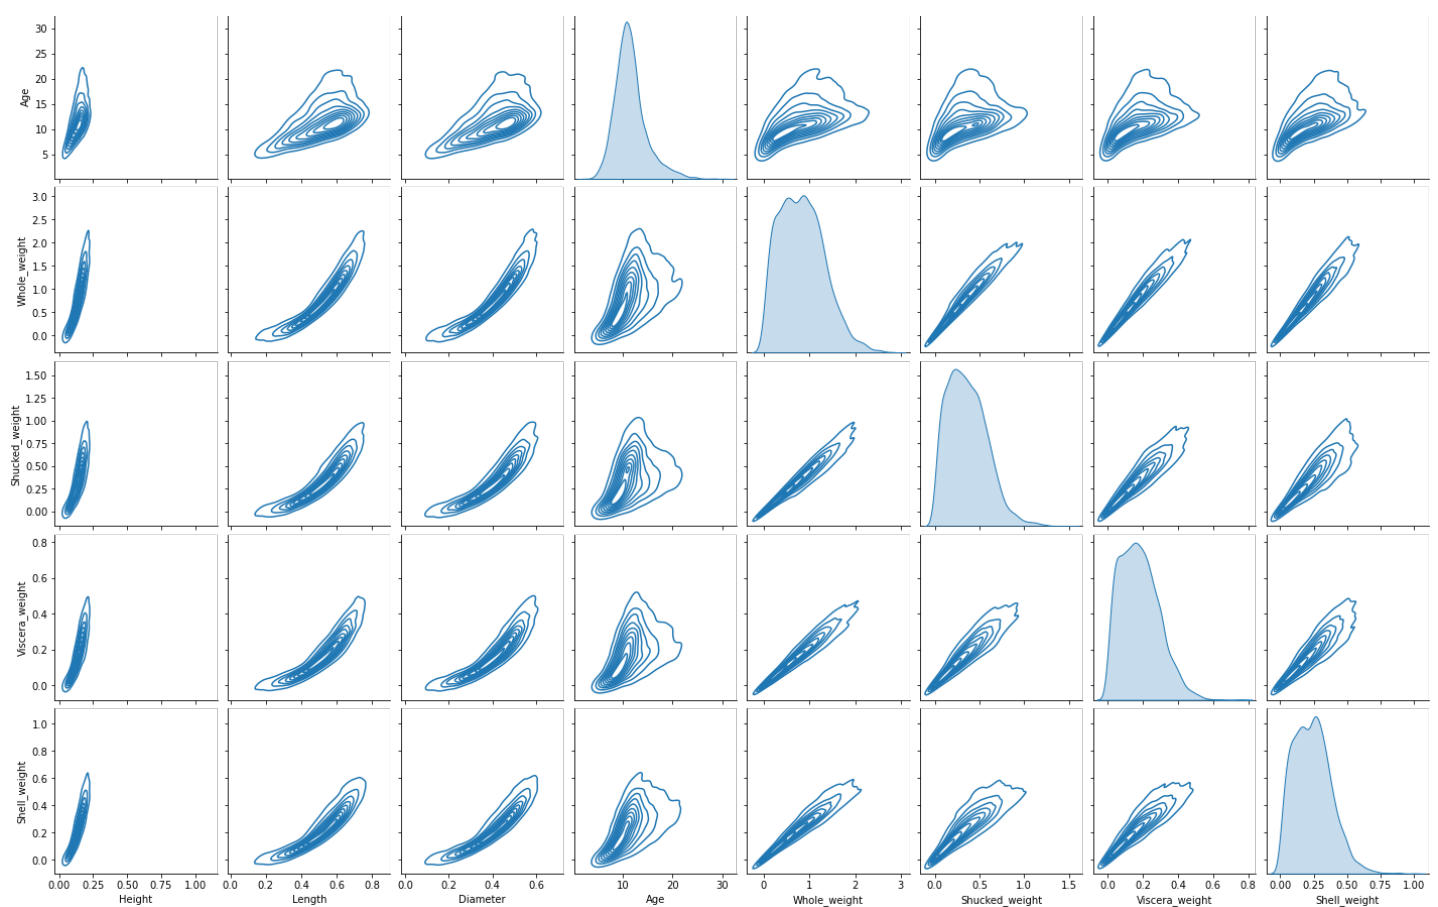
In [62]:

```
sns.pairplot(data=data[["Height", "Length", "Diameter", "Age", "Whole_weight", "Shucked_weight", "Viscera_weight", "Shell_weight"]], kind="kde")
```

Out[62]:

<seaborn.axisgrid.PairGrid at 0x7fd39840c790>





4. Perform descriptive statistics on the dataset

In [63]:

```
data.describe(include='all')
```

Out[63]:

	Sex	Length	Diameter	Height	Whole_weight	Shucked_weight	Viscera_weight	Shell_weight	Age
count	4177	4177.000000	4177.000000	4177.000000	4177.000000	4177.000000	4177.000000	4177.000000	4177.0000
unique	3	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN
top	M	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN
freq	1528	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN
mean	NaN	0.523992	0.407881	0.139516	0.828742	0.359367	0.180594	0.238831	11.4336
std	NaN	0.120093	0.099240	0.041827	0.490389	0.221963	0.109614	0.139203	3.2241
min	NaN	0.075000	0.055000	0.000000	0.002000	0.001000	0.000500	0.001500	2.5000
25%	NaN	0.450000	0.350000	0.115000	0.441500	0.186000	0.093500	0.130000	9.5000
50%	NaN	0.545000	0.425000	0.140000	0.799500	0.336000	0.171000	0.234000	10.5000
75%	NaN	0.615000	0.480000	0.165000	1.153000	0.502000	0.253000	0.329000	12.5000
max	NaN	0.815000	0.650000	1.130000	2.825500	1.488000	0.760000	1.005000	30.5000

5. Check for Missing values and deal with them

In [64]:

```
data.isnull().sum()
```

Out[64]:

Sex	0
Length	0


```

Length      0
Diameter    0
Height      0
Whole_weight 0
Shucked_weight 0
Viscera_weight 0
Shell_weight 0
Age         0
dtype: int64

```

6. Find the outliers and replace them outliers

In [65]:

```

outliers=data.quantile(q=(0.25,0.75))
outliers

```

Out[65]:

	Length	Diameter	Height	Whole_weight	Shucked_weight	Viscera_weight	Shell_weight	Age
0.25	0.450	0.35	0.115	0.4415	0.186	0.0935	0.130	9.5
0.75	0.615	0.48	0.165	1.1530	0.502	0.2530	0.329	12.5

In [66]:

```

a = data.Age.quantile(0.25)
b = data.Age.quantile(0.75)
c = b - a
lower_limit = a - 1.5 * c
data.median(numeric_only=True)

```

Out[66]:

```

Length      0.5450
Diameter    0.4250
Height      0.1400
Whole_weight 0.7995
Shucked_weight 0.3360
Viscera_weight 0.1710
Shell_weight 0.2340
Age        10.5000
dtype: float64

```

In [67]:


```

data['Age'] = np.where(data['Age'] < lower_limit, 7, data['Age'])
sns.boxplot(x=data.Age,showfliers = False)

```

Out[67]:

<AxesSubplot:xlabel='Age'>



7. Check for Categorical columns and perform encoding

In [68]:

```
data.head()
```

Out[68]:

	Sex	Length	Diameter	Height	Whole_weight	Shucked_weight	Viscera_weight	Shell_weight	Age
0	M	0.455	0.365	0.095	0.5140	0.2245	0.1010	0.150	16.5
1	M	0.350	0.265	0.090	0.2255	0.0995	0.0485	0.070	8.5
2	F	0.530	0.420	0.135	0.6770	0.2565	0.1415	0.210	10.5
3	M	0.440	0.365	0.125	0.5160	0.2155	0.1140	0.155	11.5
4	I	0.330	0.255	0.080	0.2050	0.0895	0.0395	0.055	8.5

In [83]:

```
from sklearn.preprocessing import LabelEncoder

lab = LabelEncoder()
data.Sex = lab.fit_transform(data.Sex)

data.head()
```

Out[83]:

	Sex	Length	Diameter	Height	Whole_weight	Shucked_weight	Viscera_weight	Shell_weight	Age
0	2	0.455	0.365	0.095	0.5140	0.2245	0.1010	0.150	12
1	2	0.350	0.265	0.090	0.2255	0.0995	0.0485	0.070	4
2	0	0.530	0.420	0.135	0.6770	0.2565	0.1415	0.210	6
3	2	0.440	0.365	0.125	0.5160	0.2155	0.1140	0.155	7
4	1	0.330	0.255	0.080	0.2050	0.0895	0.0395	0.055	4

8. Split the data into dependent and independent variables

In [84]:

```
y = data["Sex"]
y.head()
```

Out[84]:

```
0    2
1    2
2    0
3    2
4    1
Name: Sex, dtype: int64
```

In [85]:

```
x=data.drop(columns=["Sex"],axis=1)
x.head()
```

Out[85]:

	Length	Diameter	Height	Whole_weight	Shucked_weight	Viscera_weight	Shell_weight	Age
0	0.455	0.365	0.095	0.5140	0.2245	0.1010	0.150	12
1	0.350	0.265	0.090	0.2255	0.0995	0.0485	0.070	4
2	0.530	0.420	0.135	0.6770	0.2565	0.1415	0.210	6

3	Length	Diameter	Height	Whole_weight	Shucked_weight	Viscera_weight	Shell_weight	Age
4	0.330	0.255	0.080	0.2050	0.0895	0.0395	0.055	4

9. Scale the independent variables

In [86]:

```
from sklearn.preprocessing import scale
X_Scaled = pd.DataFrame(scale(x), columns=x.columns)
X_Scaled.head()
```

Out[86]:

	Length	Diameter	Height	Whole_weight	Shucked_weight	Viscera_weight	Shell_weight	Age
0	-0.574558	-0.432149	-1.064424	-0.641898	-0.607685	-0.726212	-0.638217	1.555152
1	-1.448986	-1.439929	-1.183978	-1.230277	-1.170910	-1.205221	-1.212987	-0.884841
2	0.050033	0.122130	-0.107991	-0.309469	-0.463500	-0.356690	-0.207139	-0.274842
3	-0.699476	-0.432149	-0.347099	-0.637819	-0.648238	-0.607600	-0.602294	0.030157
4	-1.615544	-1.540707	-1.423087	-1.272086	-1.215968	-1.287337	-1.320757	-0.884841

10. Split the data into training and testing

In [87]:

```
from sklearn.model_selection import train_test_split
X_Train, X_Test, Y_Train, Y_Test = train_test_split(X_Scaled, y, test_size=0.2, random_s
tate=0)
```

In [88]:

```
X_Train.shape, X_Test.shape
```

Out[88]:

```
((3341, 8), (836, 8))
```

In [89]:

```
Y_Train.shape, Y_Test.shape
```

Out[89]:

```
((3341,), (836,))
```

In [90]:

```
X_Train.head()
```

Out[90]:

	Length	Diameter	Height	Whole_weight	Shucked_weight	Viscera_weight	Shell_weight	Age
3141	-2.864726	-2.750043	-1.423087	-1.622870	-1.553902	-1.583867	-1.644065	-1.799838
3521	-2.573250	-2.598876	-2.020857	-1.606554	-1.551650	-1.565619	-1.626104	-1.494839
883	1.132658	1.230689	0.728888	1.145672	1.041436	0.286552	1.538726	1.555152
3627	1.590691	1.180300	1.446213	2.164373	2.661269	2.330326	1.377072	0.030157
2106	0.591345	0.474853	0.370226	0.432887	0.255175	0.272866	0.906479	1.250153

In [91]:

```
X_Test.head()
```

Out[91]:

	Length	Diameter	Height	Whole_weight	Shucked_weight	Viscera_weight	Shell_weight	Age
668	0.216591	0.172519	0.370226	0.181016	-0.368878	0.569396	0.690940	0.945154
1580	-0.199803	-0.079426	-0.466653	-0.433875	-0.443224	-0.343004	-0.325685	-0.579842
3784	0.799543	0.726798	0.370226	0.870348	0.755318	1.764639	0.565209	0.335156
463	-2.531611	-2.447709	-2.020857	-1.579022	-1.522362	-1.538247	-1.572219	-1.799838
2615	1.007740	0.928354	0.848442	1.390405	1.415417	1.778325	0.996287	0.640155

In [92]:

```
Y_Train.head()
```

Out[92]:

```
3141    1
3521    1
883     2
3627    2
2106    2
Name: Sex, dtype: int64
```

In [93]:

```
Y_Test.head()
```

Out[93]:

```
668     2
1580    1
3784    2
463     1
2615    2
Name: Sex, dtype: int64
```

11. Build the Model

In [94]:

```
from sklearn.ensemble import RandomForestClassifier
model = RandomForestClassifier(n_estimators=10,criterion='entropy')
```

In [95]:

```
model.fit(X_Train,Y_Train)
```

Out[95]:

```
RandomForestClassifier(criterion='entropy', n_estimators=10)
```

In [96]:

```
y_predict = model.predict(X_Test)
```

In [97]:

```
y_predict_train = model.predict(X_Train)
```

12. Train the Model

In [98]:

```
from sklearn.metrics import accuracy_score,confusion_matrix,classification_report
```

In [99]:

```
print('Training accuracy: ',accuracy_score(Y_Train,y_predict_train))
```

Training accuracy: 0.9787488775815624

13.Test the Model

In [100]:

```
print('Testing accuracy: ',accuracy_score(Y_Test,y_predict))
```

Testing accuracy: 0.5526315789473685

14. Measure the performance using Metrics

In [101]:

```
pd.crosstab(Y_Test,y_predict)
```

Out[101]:

col_0	0	1	2
Sex			
0	122	29	98
1	37	217	37
2	120	53	123

In [102]:

```
print(classification_report(Y_Test,y_predict))
```

	precision	recall	f1-score	support
0	0.44	0.49	0.46	249
1	0.73	0.75	0.74	291
2	0.48	0.42	0.44	296
accuracy			0.55	836
macro avg	0.55	0.55	0.55	836
weighted avg	0.55	0.55	0.55	836